

Introduction to Computer Architecture

Project 3

MIPS CPU Simulator + Cache Simulator

Hyungmin Cho

Department of Software
Sungkyunkwan University

Project 2 Overview

- In Project 3, you'll expand your project 2 implementation to model the **cache** behavior.
 - ❖ At the end of the execution, the simulator reports the number of total cache hits and misses.
 - ❖ We do not consider the instruction cache. Only the **data cache** is simulated
 - You need to model the cache behavior when you're processing the **lw** and **sw** instructions.
 - ❖ Need to support the same set of MIPS instructions as Proj2

Cache Specs to Simulate – Cache #1

- 1KB (1024 bytes) of total data capacity
- 32 bytes (8 words) per block
- Direct-mapped
- “Write-through, no write allocate” policy
- 32-bit memory address

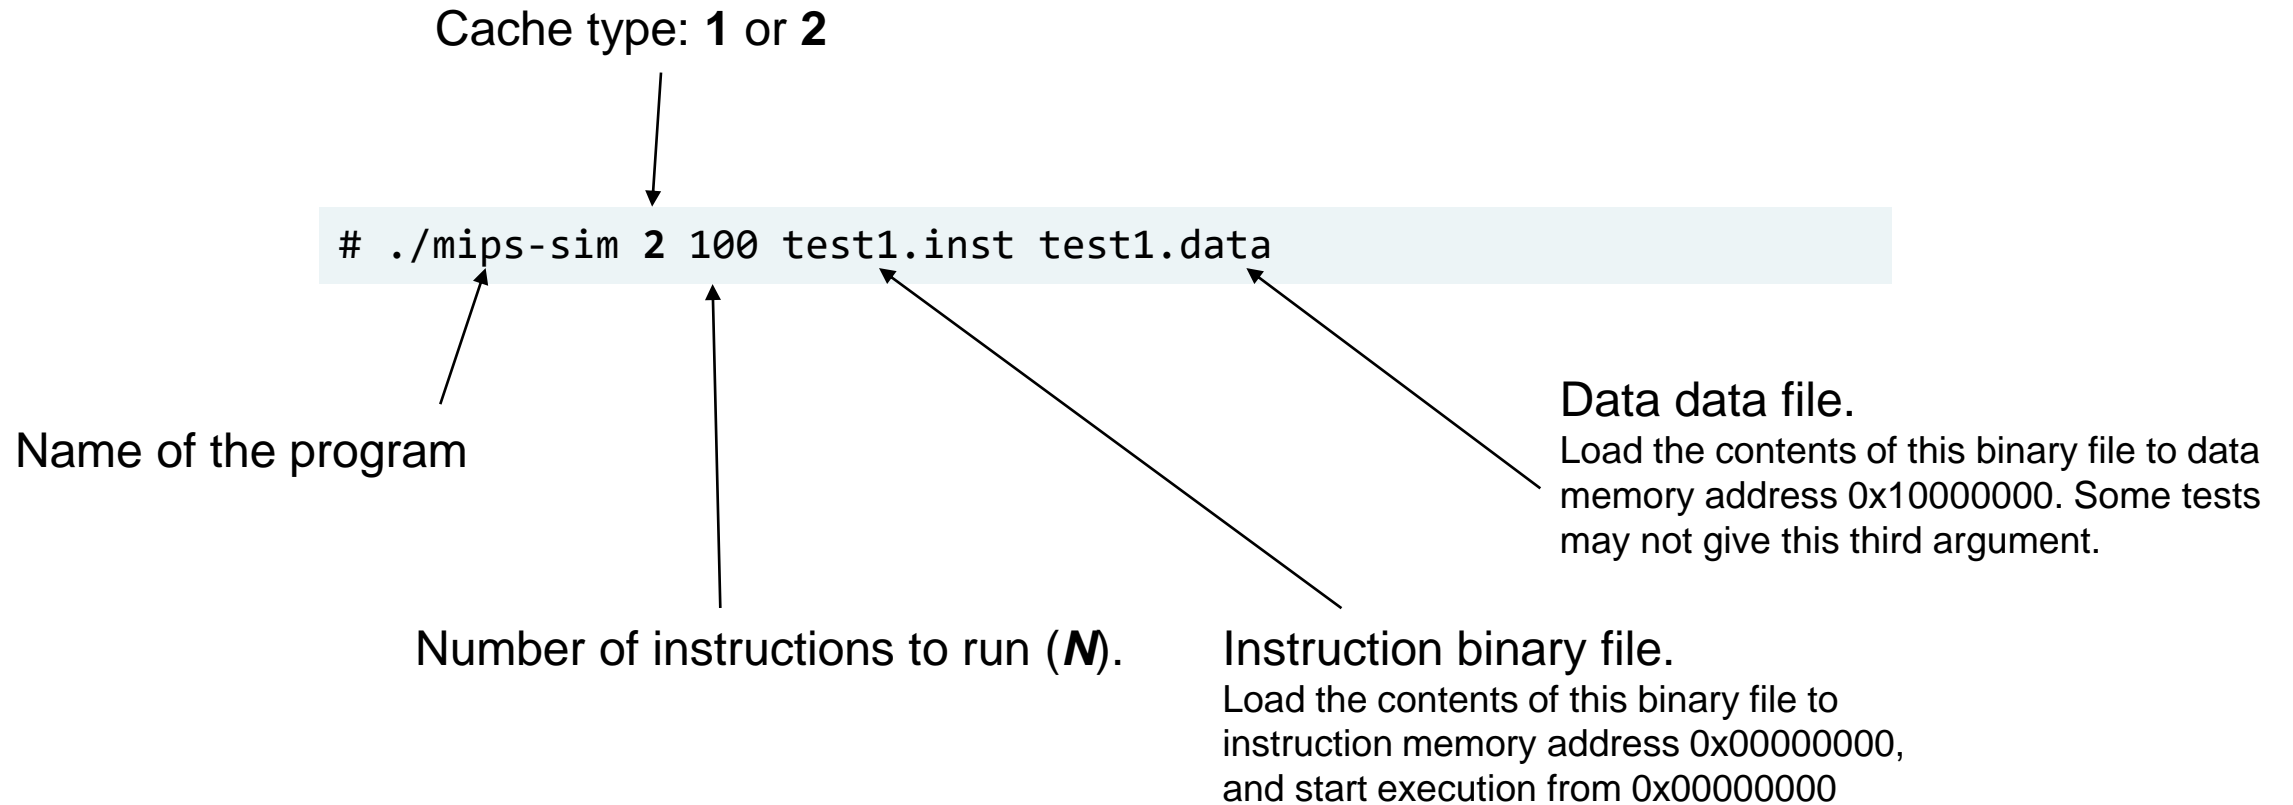
Cache Specs to Simulate – Cache #2

- 4KB (4096 bytes) of total data capacity
- 64 bytes (16 words) per block
- 4-way set associative, **True LRU** replacement policy
 - ❖ NOT the pseudo-LRU!
- “Write-back, write allocate” policy
- 32-bit memory address

Simulator Program Behavior

1. Your program takes 3 or 4 command-line arguments.
2. We add a first argument to select the cache type
 1. If the first argument is 1, simulate Cache #1
 2. If the first argument is 2, simulate Cache #2
 3. If something other than **1** or **2**, use the default value, which is **1**
3. The 2nd, 3rd, and 4th arguments are the same as the 1st, 2nd, and 3rd arguments of Proj2
 1. The second argument is the number of instructions to execute (**N**)
 2. The third argument is the instruction binary file.
 3. The fourth, optional argument is the data binary file.

Simulator Program Behavior



Outputs

1. Instructions:
 - ❖ The number of instructions executed by the simulator.
 - ❖ It can be smaller than the second argument (N) if the simulator is terminated due to an unknown instruction. In that case, print the number of executed instructions *including* the unknown instruction.
 2. Total:
 - ❖ Hits + Misses
 3. Hits:
 - ❖ The number of data cache hits
 4. Misses:
 - ❖ The number of data cache misses
-
- These numbers DO NOT include the memory accesses to load the data file. That is, start counting hits and misses after loading the data file. Also, assume all cache lines are invalid after loading the data file.
 - Do not print “unknown instruction” even if the simulator has stopped due to the unknown instruction.
 - Do not print the register values
 - Do not print the PC value

```
# ./mips-sim 1 1000 test.inst test.data
Instructions: 222
Total: 14
Hits: 10
Misses: 4
```

Reference Implementation

- We provide a reference implementation (without source code) in the following location.
 - ❖ `~swe3005/2021f/proj3/mips-sim`

Test Sample (1)

- ~swe3005/2021f/proj3/proj3_1.inst
- proj3_1 tests the cache line size difference (8 words vs. 16 words)

```
lui $2, 0x1000
lw $3, 0($2)
lw $3, 4($2)
lw $3, 8($2)
lw $3, 12($2)
...
...
lw $3, 60($2)
lw $3, 64($2)
lw $3, 68($2)
...
..
lw $3, 108($2)
lw $3, 112($2)
lw $3, 116($2)
lw $3, 120($2)
lw $3, 124($2)
nop
```

Blue ones
indicate
cache hits
for Cache #2

```
./mips-sim 1 1000 ~swe3005/2021f/proj3/proj3_1.inst
```

```
Instructions: 35
Total: 32
Hits: 28
Misses: 4
```

```
./mips-sim 2 1000 ~swe3005/2021f/proj3/proj3_1.inst
```

```
Instructions: 35
Total: 32
Hits: 30
Misses: 2
```

Test Sample (2)

- ~swe3005/2021f/proj3/proj3_2.inst
- proj3_2 tests the cache size difference (1 KB vs. 4 KB)

```
lui $2, 0x1000
add $3, $2, 0x400
lw $4, 0($2)
lw $4, 0($3)
lw $4, 4($2)
lw $4, 4($3)
nop
```

Blue ones
indicate
cache hits
for Cache #2

```
./mips-sim 1 1000 ~swe3005/2021f/proj3/proj3_2.inst
```

```
Instructions: 8
Total: 4
Hits: 0
Misses: 4
```

```
./mips-sim 2 1000 ~swe3005/2021f/proj3/proj3_2.inst
```

```
Instructions: 8
Total: 4
Hits: 2
Misses: 2
```

Test Sample (3)

- ~swe3005/2021f/proj3/proj3_3.inst
- proj3_3 tests the associativity of Cache #2 (4-way)

```
lui $2, 0x1000
add $3, $2, 0x400
add $4, $3, 0x400
add $5, $4, 0x400
add $6, $5, 0x400
lw $10, 0($2)
lw $10, 4($3)
lw $10, 8($4)
lw $10, 12($5)
lw $10, 16($2)
lw $10, 20($3)
lw $10, 24($4)
lw $10, 28($5)
lw $10, 32($2)
lw $10, 36($3)
lw $10, 40($4)
lw $10, 44($5)
lw $10, 48($6)
lw $10, 52($2)
lw $10, 56($3)
lw $10, 60($4)
lw $10, 56($5)
lw $10, 52($6)
nop
```

Blue ones
indicate
cache hits
for Cache #2

```
./mips-sim 1 1000 ~swe3005/2021f/proj3/proj3_3.inst
```

```
Instructions: 25
Total: 18
Hits: 0
Misses: 18
```

```
./mips-sim 2 1000 ~swe3005/2021f/proj3/proj3_3.inst
```

```
Instructions: 25
Total: 18
Hits: 8
Misses: 10
```

Test Sample (4)

- ~swe3005/2021f/proj3/proj3_4.inst
- proj3_4 tests the replacement scheme of Cache #2 (4-way True LRU)

```
lui $2, 0x1000
add $3, $2, 0x400
add $4, $3, 0x400
add $5, $4, 0x400
add $6, $5, 0x400
lw $10, 0($2)
lw $10, 0($3)
lw $10, 0($4)
lw $10, 0($5)
lw $10, 0($6)
lw $10, 0($3)
lui $2, 0x1000
add $2, $2, 0x40
add $3, $2, 0x400
add $4, $3, 0x400
add $5, $4, 0x400
add $6, $5, 0x400
lw $10, 0($2)
lw $10, 0($3)
lw $10, 0($4)
lw $10, 0($5)
lw $10, 0($6)
lw $10, 0($2)
lui $2, 0x1000
add $2, $2, 0x80
add $3, $2, 0x400
add $4, $3, 0x400
add $5, $4, 0x400
add $6, $5, 0x400
lw $10, 0($2)
lw $10, 0($3)
lw $10, 0($4)
lw $10, 0($5)
lw $10, 0($2)
lw $10, 0($6)
lw $10, 0($2)
lw $10, 0($3)
nop
```

Blue ones
indicate
cache hits
for Cache #2

```
./mips-sim 1 1000 ~swe3005/2021f/proj3/proj3_4.inst
```

```
Instructions: 39
Total: 20
Hits: 0
Misses: 20
```

```
./mips-sim 2 1000 ~swe3005/2021f/proj3/proj3_4.inst
```

```
Instructions: 39
Total: 20
Hits: 3
Misses: 17
```

Test Sample (5)

- ~swe3005/2021f/proj3/proj3_5.inst
- proj3_5 tests the write policy Cache #2 (Write allocate)

```
lui $2, 0x1000
addi $3, $zero, 1234
sw $3, 0($2)
lw $4, 4($2)
nop
```

Blue ones
indicate
cache hits
for Cache #2

```
./mips-sim 1 1000 ~swe3005/2021f/proj3/proj3_5.inst
```

```
Instructions: 6
Total: 2
Hits: 0
Misses: 2
```

```
./mips-sim 2 1000 ~swe3005/2021f/proj3/proj3_5.inst
```

```
Instructions: 6
Total: 2
Hits: 1
Misses: 1
```

Test Sample (6)

- ~swe3005/2021f/proj3/proj3_6.inst
- proj3_6 corresponds to the following C code

```
int main () {  
    int* p = (int*) 0x10000000;  
    int val = 0;  
  
    for(int i = 0; i < 32; i++) {  
        p[i*4] = p[i*16];  
    }  
  
    return 0;  
}
```

```
./mips-sim 1 1000 ~swe3005/2021f/proj3/proj3_6.inst
```

```
Instructions: 761  
Total: 294  
Hits: 236  
Misses: 58
```

```
./mips-sim 2 1000 ~swe3005/2021f/proj3/proj3_6.inst
```

```
Instructions: 761  
Total: 294  
Hits: 261  
Misses: 33
```

Test Sample (7)

- ~swe3005/2021f/proj3/proj3_7.inst, ~swe3005/2021f/proj3/proj3_7_8.data
- proj3_7 and proj3_8 are sorting programs. Proj3_7 implements the quicksort algorithm
- proj3_7_8.data contains 2,048 numbers (8,192 bytes)

```
./mips-sim 1 10000000 ~swe3005/2021f/proj3/proj3_7.inst ~swe3005/2021f/proj3/proj3_7_8.data
```

```
Instructions: 1137603  
Total: 434271  
Hits: 419884  
Misses: 14387
```

```
./mips-sim 2 10000000 ~swe3005/2021f/proj3/proj3_7.inst ~swe3005/2021f/proj3/proj3_7_8.data
```

```
Instructions: 1137603  
Total: 434271  
Hits: 433910  
Misses: 361
```

Test Sample (8)

- ~swe3005/2021f/proj3/proj3_8.inst, ~swe3005/2021f/proj3/proj3_7_8.data
- proj3_7 and proj3_8 are sorting programs. Proj3_8 implements the radix sort algorithm
- proj3_7_8.data contains 2,048 numbers (8,192 bytes)
- *If you look at the results closely, proj3_8 executes fewer instructions, but has a higher number of cache misses.*

```
./mips-sim 1 10000000 ~swe3005/2021f/proj3/proj3_8.inst ~swe3005/2021f/proj3/proj3_7_8.data
```

```
Instructions: 1067581  
Total: 402343  
Hits: 378040  
Misses: 24303
```

```
./mips-sim 2 10000000 ~swe3005/2021f/proj3/proj3_8.inst ~swe3005/2021f/proj3/proj3_7_8.data
```

```
Instructions: 1067581  
Total: 402343  
Hits: 400123  
Misses: 2220
```


Project Environment

- We will use the department's In-Ui-Ye-Ji cluster
 - ❖ `swui.skku.edu`
 - ❖ `swye.skku.edu`
 - ❖ `swji.skku.edu`
 - ❖ ssh port: 1398
- You'll need a similar Makefile as proj2
 - ❖ Same executable file name (i.e., `mips-sim`)
- If you have a problem with the account, send an e-mail to the server admin
 - ❖ inuiyeji-skku@googlegroups.com
 - ❖ Do not send an email that is not related to the account itself!
 - ❖ If you're not sure, ask the TAs first.

Submission

- Clear the build directory
 - ❖ Do not leave any executable or object file in the submission
 - ❖ `make clean`
- Use the submit program
 - ❖ `~swe3005/bin/submit project_id path_to_submit`
 - ❖ If you want to submit the 'project_3' directory...
 - `~swe3005/bin/submit proj3 project_3`

Submitted Files for proj2:

File Name	File Size	Time

proj3-2020123456-Sep.05.17.22.388048074	268490	Thu Sep 5 17:22:49 2020

- Verify the submission
 - ❖ `~swe3005/bin/check-submission proj3`

Project 2 Due Date

- 2021 Dec 13th, 23:59:59
- No late submission